



Preparing images for deep learning

Normalization / Standardization / Feature Scaling

- Standardization (most widely used)
 - Works well for populations that are normally distributed
 - Output can be -ve (bad for methods that expect positive inputs, e.g. RBMs)

$$x' = \frac{x - \bar{x}}{\sigma}$$

- Rescaling (min-max normalization)
 - Very fast

$$x' = \frac{x - \min(x)}{\max(x) - \min(x)}$$

- Mean normalization

$$x' = \frac{x - \text{average}(x)}{\max(x) - \min(x)}$$

- Rescaling by dividing by maximum (most common for images)
- Why data normalization or standardization?
 - In stochastic gradient descent, feature scaling can improve the convergence speed of the algorithm
 - In support vector machines, it can reduce the time to find support vectors

5.2.4 Data preprocessing

- We should load the data into memory to do the training
 - Here is what we have been doing:

In future, we will have X but not Y!

Input (X)							Output (Y)
Fixed acidity	Volatile acidity	Citric acid	Residual sugar	...	Sulphates	Alcohol	Quality
7.4	0.70	0.00	1.90		0.6	9.4	5
7.8	0.88	0.00	2.60		0.7	9.8	5
7.8	0.76	0.04	2.30		0.7	9.8	5
11.2	0.28	0.56	1.90		0.6	9.8	6
7.4	0.70	0.00	1.90		0.6	9.4	5
...
7.9	0.60	0.06	1.60		0.5	9.4	5
7.3	0.65	0.00	1.20		0.5	10.0	7
7.8	0.58	0.02	2.00		0.6	9.5	7
7.5	0.50	0.36	0.10		0.8	10.5	5
6.7	0.58	0.08	1.80		0.5	9.2	5

Split for training and validation

Fixed acidity	Volatile acidity	Citric acid	Residual sugar	...	Sulphates	Alcohol	Quality
7.4	0.70	0.00	1.90		0.6	9.4	5
7.8	0.88	0.00	2.60		0.7	9.8	5
7.8	0.76	0.04	2.30		0.7	9.8	5
11.2	0.28	0.56	1.90		0.6	9.8	6
7.4	0.70	0.00	1.90		0.6	9.4	5
...
7.9	0.60	0.06	1.60		0.5	9.4	5
7.3	0.65	0.00	1.20		0.5	10.0	7
7.8	0.58	0.02	2.00		0.6	9.5	7
7.5	0.50	0.36	0.10		0.8	10.5	5
6.7	0.58	0.08	1.80		0.5	9.2	5

Available data

XTRAIN

YTRAIN

XVALID

YVALID

```
history = model.fit(XTRAIN, YTRAIN, validation_data=(XVALIDATION, YVALIDATION), epochs=256)
```

- If we have Terabytes of data, can we load the entire dataset into XTRAIN and YTRAIN? Why?

An ideal way of loading the data & training

- As the training starts:
 1. Read a picture file
 2. Decode the JPEG content to RGB grids of pixels
 3. Convert these into floating-point tensors
 4. Rescale the pixel values (between 0 and 255) to the $[0, 1]$ interval
 5. Add it to the pool of training/validation dataset
- Keras has utilities to take care of these steps automatically
 - contains the class **ImageDataGenerator** which lets you quickly set up Python generators that can automatically turn image files on disk into batches of preprocessed tensors

Python generators

- Python has a 'yield' operator that you can use in place of 'return'
- A Python generator is an object that acts as an iterator
 - You can use it with the for ... in operator

```
1 def generator():  
2     i = 0  
3     while True:  
4         i += 1  
5         yield i
```

```
1 for x in generator():  
2     print(x)  
3     if x > 5:  
4         break
```

5.7 Data generator for generating images

```
from keras.preprocessing.image import ImageDataGenerator
```

```
train_datagen = ImageDataGenerator(rescale=1./255)
```

Rescales all images by 1/255

```
train_generator = train_datagen.flow_from_directory(
```

```
    train_dir,
```

```
    target_size=(150, 150)
```

Resizes all images to 150 × 150

```
    batch_size=20,
```

```
    class_mode='binary')
```

**Because you use
binary_crossentropy
loss, you need binary
labels.**

```
history = model.fit_generator(
    train_generator,
    steps_per_epoch=100,
    epochs=10,
    validation_data=validation_generator,
    validation_steps=50)
```

**Target
directory**

5.2.5 Data augmentation

- The approach of generating more training data from existing training samples
 - by augmenting the samples via a number of random transformations that yield believable-looking images
- The goal is: “at training time, the model will never see the exact same picture twice”
 - This helps expose the model to more aspects of the data and generalize better
- If we read images using `ImageDataGenerator` instance (in Keras)
 - a number of random transformations can be performed

```
from keras.preprocessing.image import ImageDataGenerator

train_datagen = ImageDataGenerator(rescale=1./255)
test_datagen = ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(150, 150)
    batch_size=20,
    class_mode='binary')

# Resizes all images to 150 x 150
# Because you use binary_crossentropy loss, you need binary labels.
```

Target directory

Previous Code

```
datagen = ImageDataGenerator(
    rotation_range=40,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest')
```

Augmentation using ImageDataGenerator

- **rotation_range**
a value in degrees (0–180), a range within which to randomly rotate pictures
- **width_shift** and **height_shift**
ranges (as a fraction of total width or height) within which to randomly translate pictures vertically or horizontally
- **shear_range**
randomly applying shearing transformations.
- **zoom_range**
randomly zooming inside pictures.
- **horizontal_flip**
is for randomly flipping half the images horizontally—relevant when there are no assumptions of horizontal asymmetry (for example, real-world pictures)
- **fill_mode**
the strategy used for filling in newly created pixels, which can appear after a rotation or a width/height shift

... (there are more)

```
datagen = ImageDataGenerator(  
    rotation_range=40,  
    width_shift_range=0.2,  
    height_shift_range=0.2,  
    shear_range=0.2,  
    zoom_range=0.2,  
    horizontal_flip=True,  
    fill_mode='nearest')
```