



# From data to a deep learning model: Deep learning workflow / Recipe

# The universal workflow of machine learning

Step 1: Define the problem and assemble a dataset ✓

- Binary / multi-class classification or regression?

Step 2: Choose a measure of success ✓

- Accuracy (for balanced), Precision / Recall (for unbalanced)

Step 3: Decide on an evaluation protocol ✓

- Holdout / K-fold cross validation

Step 4: Prepare your data ✓

- Normalization / standardization

Step 5: Develop a model that does better than a baseline ✓

Step 6: Scale up: Develop a model that overfits ✓

Step 7: Regularize your model and tune hyperparameters ✓

## Step 5: Develop a model that does better than the baseline

- Can you achieve “statistical power”?
  - a small model that is capable of beating a dumb baseline
- Example 1:
  - In the MNIST digit-classification example, anything that achieves an accuracy greater than 0.1
- Example 2:
  - In the IMDB example, it's anything with an accuracy greater than 0.5

## Step 5: Develop a model that does better than the baseline

- Can you achieve “statistical power”?
  - a small model that is capable of beating a dumb baseline
- Example 1:
  - In the MNIST digit-classification example, anything that achieves an accuracy greater than 0.1
- Example 2:
  - In the IMDB example, it's anything with an accuracy greater than 0.5
- What can you do, if you cannot build a model with statistical power?
  - Check the last-layer's activation - are you using sigmoid for regression?
  - Check the loss function - `binary_crossentropy` for classification and `mse/mae` for regression
  - Check optimizer - start with `rmsprop` with its default learning rate
  - Last resort: Use output as one of the input features

## Step 6: Scale up: Develop a model that overfits

- So, we obtained a model that has statistical power, what next?
- Is our model sufficiently powerful to learn more patterns?
  - Does it have enough layers and parameters to properly model the problem at hand?

# Step 6: Scale up: Develop a model that overfits

- So, we obtained a model that has statistical power, what next?
- Is our model sufficiently powerful to learn more patterns?
  - Does it have enough layers and parameters to properly model the problem at hand?
- To figure out if you can overfit, you have to overfit
  - How to overfit?
    - Add layers, Make the layers bigger, and Train for more epochs
  - Monitor the training loss and validation loss (and accuracy/mae)
  - When you see that the model's performance on the validation data begins to degrade, you've achieved overfitting

## Step 7: Regularize the model and tune hyper-parameters

- Repeatedly modify your model, train it, evaluate on your validation data
  - Repeat, until the model is as good as it can get
- What to change?
  - Add dropout
  - Try different architectures: add or remove layers
  - Add L1 and/or L2 regularization
  - Try different hyperparameters (such as the number of units per layer or the learning rate of the optimizer) to find the optimal configuration
  - [Optional] Add new features or remove features that don't seem to be informative

## Step 7: Regularize the model and tune hyper-parameters

- Repeatedly modify your model, train it, evaluate on your validation data
  - Repeat, until the model is as good as it can get
- What to change?
  - Add dropout
  - Try different architectures: add or remove layers
  - Add L1 and/or L2 regularization
  - Try different hyperparameters (such as the number of units per layer or the learning rate of the optimizer) to find the optimal configuration
  - [Optional] Add new features or remove features that don't seem to be informative
- Once you've developed a satisfactory model configuration
  - Train your final production model on all the available data (training and validation) and evaluate it one last time on the test set



# How to debug a deep learning development pipeline?

