




Overfitting, underfitting and regularization (Bias-variance tradeoff)

Underfitting
(high-bias)




Overfitting
(high-variance)

Overfitting (variance) and underfitting (bias)

“The bias-variance tradeoff is a central problem in supervised learning”

Underfitting
(high-bias)



Overfitting
(high-variance)

High bias is too less optimization (having too few blind people)

- Model misses to learn many representative training data


High variance is too much optimization (having too many blind people)

- Model learns unrepresentative training data as well

We do not want either!!

We want generalization:

- A few but ‘smart’ blind person “One side is snake like and the other side rope like” or “...”, etc.



How to prevent overfitting?

We fight overfitting using “regularization” techniques:

1. Get more training data (more real data)
2. Data augmentation (more fake data)
3. Reduce the network's size
 - Number of neurons and layers
4. Add weight regularization
 - L1 and L2 penalties
5. Adding dropout layers
6. Early stopping and model checkpointing
7. Adding batch normalization layers (not fully considered a regularization technique)

3. Regularization by reducing the network's size

- Model with more parameters = More memorization capacity
 - More capacity allows a model to easily learn a perfect dictionary-like mapping between 'training samples' and their 'targets'
 - a mapping 'without any generalization power'
- The challenge is generalization, not fitting
 - Deep learning models tend to be good at fitting to the training data
- "Too much capacity vs not enough capacity"
 - If the network has limited memorization resources, it won't be able to learn this mapping as easily
 - The model will starve for memorization resources
- So, what is the magical formula?
 - We must evaluate an array of different architectures (on your validation set, not on test set) in order to find the correct model size for your data
 - The general workflow to find an appropriate model size:
 - Start with relatively few layers and parameters, and increase the size of the layers or add new layers until you see diminishing returns with regard to validation loss

How do I know if my model is already too large?

Listing 4.3 Original model

```
from keras import models
from keras import layers



model = models.Sequential()
model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

Listing 4.4 Version of the model with lower capacity

```
model = models.Sequential()
model.add(layers.Dense(4, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(4, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

Listing 4.5 Version of the model with higher capacity

```
model = models.Sequential()
model.add(layers.Dense(512, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```



Which network nears zero loss more quickly?

4. Adding weight regularization

- **Occam's razor:** Given two explanations for something, the explanation most likely to be correct is the simplest one—the one that makes fewer assumptions
- Simpler models are less likely to overfit than complex ones
 - A simple model in this context is a model where the distribution of parameter values has less entropy (or a model with fewer parameters)
- How to reduce the complexity of a network (without changing # of parameters?)
 - Force its weights to take only small values
 - This will make the distribution of weight values more **regular**
- Weight regularization
 - Adding to the loss function of the network - “a cost associated with having large weights”
 - I.e. don't let the weight values grow

Two ways to regularize weights: L_1 and L_2

1. L_1 regularization

- The cost added is proportional to the absolute value of the weight coefficients

2. L_2 regularization


- The cost added is proportional to the square of the value of the weight coefficients
- L_2 regularization is also called weight decay in the context of neural networks

```
from tensorflow.keras import regularizers  
  
...  
model.add(Conv2D(4,(3,3), kernel_regularizer=regularizers.l1, ...))  
model.add(Flatten())  
model.add(Dense(10, kernel_regularizer=regularizers.l2(0.001), ...))  
...
```

What does `l2(0.001)` do?

Every coefficient in the weight matrix of the layer will add

$0.001 * \text{weight_coefficient_value}$ to the total loss of the network.



5. Adding dropout

- Dropout is one of the most effective and most commonly used regularization
- Dropout, applied to a layer, consists of randomly dropping out (setting to zero) a number of output features of the layer during training
 - Let's say a given layer would normally return a vector $[0.2, 0.5, 1.3, 0.8, 1.1]$ for a given input sample during training
 - After applying dropout, this vector will have a few zero entries distributed at random
 - For example, $[0, 0.5, 1.3, 0, 1.1]$
- The **dropout rate** is the fraction of the features that are zeroed out
 - It's usually set between 0.2 and 0.5

```
...
model.add(Conv2D(16, 3, activation = 'sigmoid', ...))
model.add(Dropout(rate = 0.25))
model.add( Conv2D( 4, ( 3, 3 ), activation = 'relu' ) )
...
```



How is dropout handled during testing?

Strategy 1: At test time, no units are dropped out

- Instead, the layer's output values are scaled down by a factor equal to the dropout rate
- `layer_output *= 0.5`
- This is to balance for the fact that more units are active than at training time

Strategy 2: Inverted dropout

- `activation = activation / keep_probability` to scale up during the training



Why does dropout work?

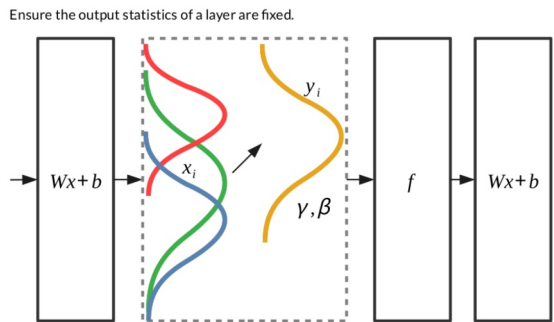
- Hinton was inspired by a fraud-prevention mechanism used by banks
 - “I went to my bank. The tellers kept changing and I asked one of them why. He said he didn’t know but they got moved around a lot. I figured it must be because it would require cooperation between employees to successfully defraud the bank. This made me realize that randomly removing a different subset of neurons on each example would prevent conspiracies and thus reduce overfitting.”
- The core idea is that introducing noise in the output values of a layer can break up **happenstance patterns** that aren’t significant

7. Batch normalization

Problem: The distribution of each layer's input changes during training

Solution: Fix the distribution

Batch normalization is a key solution to avoid/minimize the vanishing gradient problem



```
model.add( Conv2D( 16, ( 3, 3 ), activation = 'sigmoid', input_shape = xtrain[0, :, :, :].shape ) )  
model.add( BatchNormalization() )  
model.add( Conv2D( 4, ( 3, 3 ), activation = 'relu' ) )
```

Example

..**AUGMENTED INPUT**..

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Conv2D, Flatten, BatchNormalization, MaxPool2D, Dropout
model = Sequential()
model.add( Conv2D( 64, ( 3, 3 ), activation = 'sigmoid', input_shape = xtrain[0, :, :, :].shape ) )
model.add(BatchNormalization())
model.add( Conv2D( 64, ( 3, 3 ), activation = 'relu' ) )
model.add(MaxPool2D(2, 2))
model.add( Conv2D( 64, ( 3, 3 ), activation = 'relu' ) )
model.add(Dropout(0.25))
model.add( Conv2D( 64, ( 3, 3 ), activation = 'relu' ) )
model.add(BatchNormalization())
model.add( Conv2D( 64, ( 3, 3 ), activation = 'relu' ) )
model.add(BatchNormalization())
model.add( Conv2D( 64, ( 3, 3 ), activation = 'relu' ) )
model.add(BatchNormalization())
model.add( Flatten() )
model.add( Dense( 10, activation = 'relu' ) )
model.add( Dense( 10, activation = 'softmax' ) )
model.summary()
```

..**TRAINING WITH CALLBACKS**..

Training a very deep ConvNet: The villains

```
model = Sequential()
model.add(Conv2D(filters = 16, kernel_size = 5, activation = 'relu', input_shape = (28,28,1)))

for i in range(1000):
    model.add(Conv2D(filters = 4, kernel_size = 5, activation = 'relu', padding='same'))

model.add(Conv2D(filters = 16, kernel_size = 5, activation = 'relu'))
model.add(MaxPooling2D(pool_size = 2, strides = 2))
model.add(Flatten())
model.add(Dense(units = 16, activation = 'relu'))
model.add(Dense(units = 16, activation = 'relu'))
model.add(Dense(units = 10, activation = 'softmax'))
model.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metrics = ['accuracy'])

print(model.summary())
```

Some major villains when training ‘very deep’ networks:

- Vanishing Gradient, Exploding Gradient, and Internal Covariate Shift

‘ReLU’ and ‘BatchNormalization’ greatly resolve this

- But, BatchNormalization slows down training